Steps:

1. Create HDInsight resource
2. Upload Notebook and Walmart.csv file in Azure Storage
3. Complete exercises on Notebook

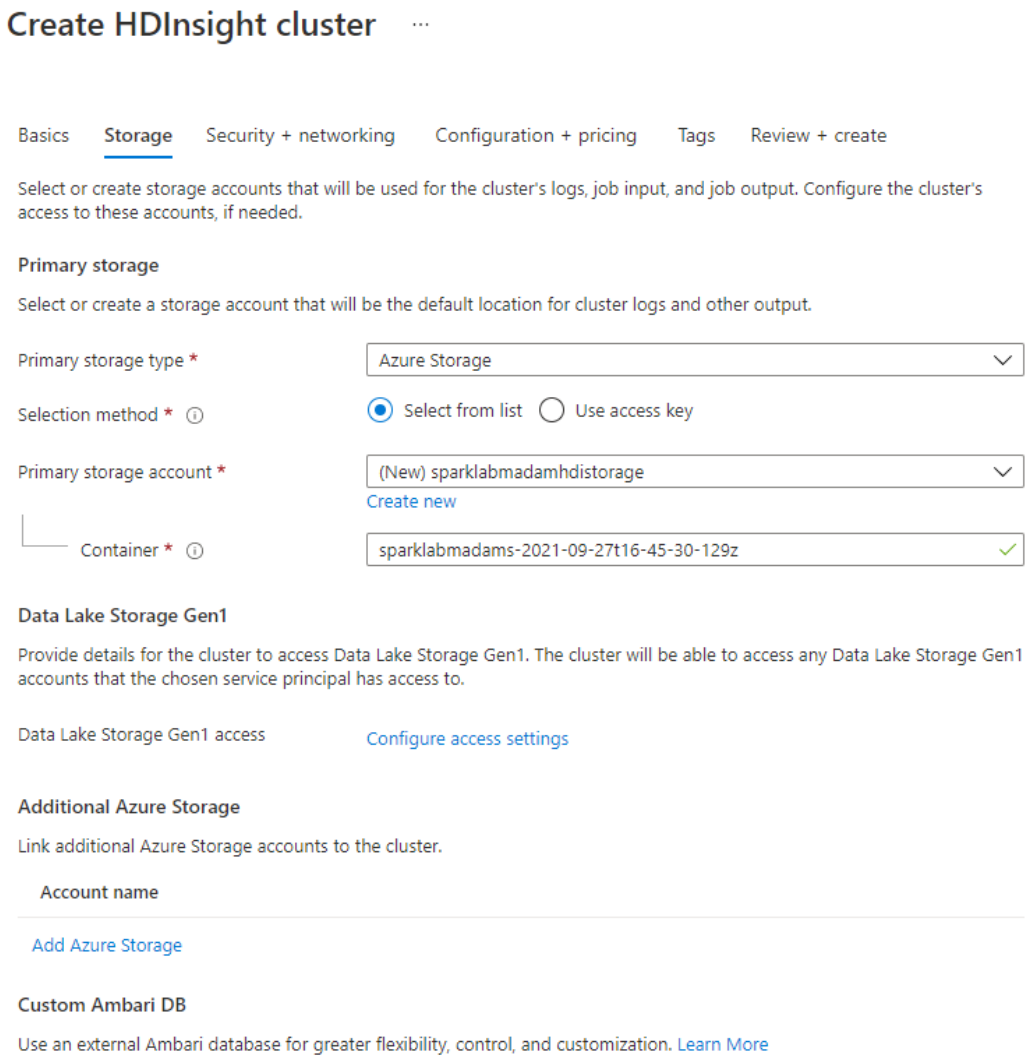

Create HDInsight cluster ...

TLS

Select the minimum TLS version supported for your cluster. Learn More

Minimum TLS version ⓘ: 1.2

Network settings

Resource provider connection ⓘ: Inbound

Connect this cluster to a virtual network. Learn More

Virtual network ⓘ:

Encryption in transit

Configure encryption in transit settings. Learn More

☐ Enable encryption in transit ⓘ

Encryption at rest

Configure disk encryption settings. Learn More

☐ Provide your own key from key vault ⓘ

☐ Enable encryption at host on temp data disk ⓘ

Identity

Select a user-assigned service identity to represent your cluster for enterprise security package or disk encryption. Learn More

User-assigned managed identity ⓘ:

Review + create | « Previous | Next: Configuration + pricing »

Create HDInsight cluster ...

Basics | Storage | Security + networking | **Configuration + pricing** | Tags | Review + create

Configure cluster performance and pricing. Learn More

Node configuration

Configure your cluster's size and performance, and view estimated cost information.

The cost estimate represented in the table does not include subscription discounts or costs related to storage, networking, or data transfer.

There are not enough cores available to support the selected number of nodes. Please adjust the number of nodes selected, pick a different region, or open a support case to request additional HDInsight cores.
View cores usage
Open an HDInsight quota increase support case

+ Add application

Node type	Node size	Number of ...	Estimated cost/h...
Head node	D14 v2 (16 Cores, 112 GB RAM), 1.50 USD/hour	2	2.99 USD
Zookeeper node	A2 v2 (2 Cores, 4 GB RAM), 0.13 USD/hour	3	0.00 (FREE)
Worker node	D12 v2 (4 Cores, 28 GB RAM), 0.37 USD/hour	4	1.50 USD

You have reached your subscription's cores quota limit in East US. Please choose a different region or request billing support to increase your limit for East US. The value must be between 1 and 1.

☐ Enable autoscale
Learn More

Total estimated cost/hour **4.49 USD**

Script actions

Use script actions to run custom PowerShell or Bash scripts on cluster nodes during cluster provisioning. Learn about script actions

+ Add script action

Create HDInsight cluster ...

Basics | Storage | Security + networking | Configuration + pricing | **Tags** | Review + create

Name ⓘ		Value ⓘ	Resource
	:		HDInsight cluster

Azure Blob Storage on IoT Edge

Microsoft

Azure Blob Storage on IoT Edge

Add to Favorites

Microsoft

Create

Overview | Plans | Usage Information + Support | Reviews

Azure Blob Storage module on IoT Edge enables Edge-local applications that use Azure Storage SDK to alternatively store the data locally on an Edge-local blob store. This module allows you to configure it to automatically upload the data from edge to Azure, and provides support for intermittent internet connectivity. It also allows you to configure it to automatically delete the data on edge.

Having edge-local blob storage would be invaluable in scenarios that require low latency access for local processing, e.g. local image or video processing before the processed data can be transferred to the cloud at infrequent intervals. The blob store location can be switched from the cloud to the edge with a simple change of connection string in code. Further, having Azure-consistent blob storage support on IoT Edge enables business logic of an app written to run on public cloud storage to be re-hosted on to the edge any time, while requiring no code changes in the application itself.

Azure Blob Storage on IoT Edge is a module designed for lightweight multi-OS deployment on top of the IoT Edge platform. This module places the edge-local blob store on